


AI Jury System

AI-Powered Hackathon Evaluation System

Evaluation Report

Generated: 8/8/2026, 4:40:39 PM



This report belongs to:
Team 4

Team	Team 4
Language	EN
Total Score	293 / 377
Grade	Good

CATEGORY SCORES

Problem Statement 22/30 73%	Backend 41/50 82%
Frontend 32/40 80%	LLM Usage 45/60 75%
AI Engineering 49/60 82%	Security & HITL 26/40 65%
MVP & Outcome 44/50 88%	Presentation & Innovation 25/37 68%

Innovation & Business Value
This report is auto-generated by AI Jury System. For official results contact the evaluation panel.

Executive Summary

293 / 377 points

77.7% — Good

CATEGORY-WISE PERFORMANCE

Problem Statement	22/30		73%	Good
Backend	41/50		82%	Good
Frontend	32/40		80%	Good
LLM Usage	45/60		75%	Good
AI Engineering	49/60		82%	Good
Security & HITL	26/40		65%	Average
MVP & Outcome	44/50		88%	Excellent
Presentation & Innovation	25/37		68%	Average
Innovation & Business	9/10		90%	Excellent

STRENGTHS & WEAKNESSES SUMMARY

Key Strengths

- Strong semantic match to the logistics challenge through multi-agent traffic/weather/fleet analysis, ETA calculation, and explainable recommendations.
- Useful end-user surfaces exist for planners and tracking users, including map visualization, alerts, dashboards, and human-in-the-loop recommendation handling.
- Backend architecture strongly matches the PPT and success checklist with orchestrator, specialized agents, decision engine, governance, and logistics-facing APIs.
- Configuration/security hygiene is solid for a hackathon prototype, with environment-based settings and no detected hardcoded secrets.
- Broad coverage of required logistics user journeys across planner dashboard, route optimization, AI recommendations, analytics, and public shipment tracking.
- Good evidence of user feedback and graceful degradation via alerting, polling resilience, and deterministic fallback behavior.

Areas for Improvement

- Quantitative success checklist items such as 20% delivery-time reduction and sub-3 second recalculation latency are not evidenced.
- Automated integration/load testing and explicit privacy/compliance implementation are not clearly shown in the provided code evidence.
- Operational persistence is inconsistent because the app mixes SQLAlchemy schema design with in-memory demo state for many live flows.
- Error handling and separation of concerns are uneven across endpoints, with some routes returning ad hoc error JSON and main.py carrying too much logic.
- Limited direct evidence of explicit loading states and robust accessibility implementation.
- Static analysis flagged at least one unmatched frontend API call (`/health`), so not all flows are fully validated end-to-end.
-

Category Score Overview

Category	Score	Max	Distribution	%	Conf.	Grade
Problem Statement PS	22	30		73%	Medium	Good
Backend BACK	41	50		82%	High	Good
Frontend FRONT	32	40		80%	Medium	Good
LLM Usage LLM	45	60		75%	Medium	Good
AI Engineering AI_ENG	49	60		82%	Medium	Good
Security & HITL SECURITY_HITL	26	40		65%	Medium	Average
MVP & Outcome MVP	44	50		88%	High	Excellent
Presentation & Innovation PRESENTATION	25	37		68%	Medium	Average
Innovation & Business Value INNOVATION	9	10		90%	Medium	Excellent

Detailed Category Analysis

Problem Statement (PS)

22 / 30

73% — Good

The submission aligns well with the problem statement's core intent: a generative-AI-assisted multi-agent logistics platform with route intelligence, ETA updates, tracking, and planner/customer-facing visibility. I relied primarily on the problem statement plus retrieved backend/frontend code and the capability map, especially because the Integration Map explicitly said backend route detection was unavailable for this stack and should not be treated as negative evidence. The main scoring limitation is not lack of product functionality, but lack of hard proof for mandatory quantitative success metrics, compliance controls, and automated testing.

Strengths

- Strong semantic match to the logistics challenge through multi-agent traffic/weather/fleet analysis, ETA calculation, and explainable recommendations.
- Useful end-user surfaces exist for planners and tracking users, including map visualization, alerts, dashboards, and human-in-the-loop recommendation handling.

Areas to Improve

- Quantitative success checklist items such as 20% delivery-time reduction and sub-3 second recalculation latency are not evidenced.
- Automated integration/load testing and explicit privacy/compliance implementation are not clearly shown in the provided code evidence.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	Core business objective of the problem statement is clearly achieved.	PARTIAL	5/6
Evidence: frontend/src/components/map/GoogleMapComponent.tsx: "Renders vehicles, shipment origin/destination/current markers, route polylines, and hub locations on a live Google Map"; backend/app/services/orchestrator.py: computes...			
2	All mandatory features listed in the problem statement are implemented.	PARTIAL	4/6
Evidence: backend/app/agents/traffic_agent.py and backend/app/agents/weather_agent.py implement specialized agents using live tool calls plus Gemini reasoning; frontend/src/components/map/GoogleMapComponent.tsx and Capability Map ...			
3	Solution produces correct / useful outputs for the intended use case.	PARTIAL	5/6
Evidence: backend/app/services/orchestrator.py: uses route availability and aggregated delay factors to compute dynamic ETA; backend/app/services/llm_service.py prompt instructs model to "Do NOT invent facts" and "Do NOT change ca...			
4	Edge cases relevant to the problem statement are reasonably handled.	PARTIAL	5/6
Evidence: backend/app/agents/traffic_agent.py: "Falls back to a deterministic scoring formula if no GEMINI_API_KEY is set or the call fails"; backend/app/agents/weather_agent.py has the same fallback behavior; frontend/src/utlis/g...			
5	Success criteria implied by the problem statement are met.	PARTIAL	3/6
Evidence: PPT Slide 1 and Slide 3 show a "STABLE PROTOTYPE" with "6 AGENTS ACTIVE," route optimization, freight tracking, and human-in-the-loop approval; backend/app/services/llm_service.py and governance/decision-engine reference...			

% — Good

The backend is a credible FastAPI-based logistics prototype that aligns well with the problem statement: it includes multi-agent orchestration, simulation, routing/weather integrations, shipment tracking, and governance/audit concepts. I relied primarily on the problem checklist and retrieved code because the Integration Map explicitly said backend route detection was unavailable for this stack, so I did not treat that as a negative. Scores were reduced mainly where production-grade persistence, validation consistency, and architectural separation were only partial despite strong prototype functionality.

Strengths

- Backend architecture strongly matches the PPT and success checklist with orchestrator, specialized agents, decision engine, governance, and logistics-facing APIs.
- Configuration/security hygiene is solid for a hackathon prototype, with environment-based settings and no detected hardcoded secrets.

Areas to Improve

- Operational persistence is inconsistent because the app mixes SQLAlchemy schema design with in-memory demo state for many live flows.
- Error handling and separation of concerns are uneven across endpoints, with some routes returning ad hoc error JSON and main.py carrying too much logic.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	API routes / backend endpoints are logically structured. <i>Evidence: backend/app/api/simulation.py defines an APIRouter with prefix="/api/simulation"; backend/app/routes/recommendations.py defines prefix="/api/recommendations"; retrieved main.py snippets show grouped endpoints such as /ap...</i>	MET	8/9
2	Data persistence (DB, file storage, cache) is implemented appropriately. <i>Evidence: backend/app/database.py sets up SQLAlchemy engine/session using settings.DATABASE_URL; backend/app/models/shipment.py, route.py, audit.py define persistent ORM tables; backend/app/seeds/seed_data.py seeds Shipment/Vehicl...</i>	PARTIAL	7/9
3	Error handling and input validation exist on backend endpoints. <i>Evidence: backend/app/api/simulation.py accepts a typed SimulationRequest and wraps execution in try/except, converting ValueError to HTTP 404 and generic exceptions to HTTP 500. Retrieved snippets from backend/app/routes/shipment...</i>	PARTIAL	5/8
4	Backend architecture supports the scale/scope described in the PPT. <i>Evidence: PPT slide 3-4 describes multi-agent intelligence, decision engine, governance, human approval, and command center. Code evidence includes backend/app/orchestrator/orchestrator.py and app/services/orchestrator.py for orch...</i>	MET	7/8
5	Business logic is separated from I/O / presentation concerns. <i>Evidence: Business logic is separated into backend/app/engines/decision_engine.py, backend/app/engines/governance.py, backend/app/services/governance_service.py, backend/app/services/orchestrator.py, and backend/app/agents/*.py. H...</i>	PARTIAL	6/8

6 Configuration and secrets are not hardcoded insecurely.

MET

8/8

Evidence: backend/app/config.py uses Pydantic settings for DATABASE_URL, CORS, and API keys; BE layer notes state values load from environment variables or .env. STATIC EVIDENCE reports 'No hardcoded secrets/API keys in source: PA...

Frontend (FRONT)

32 / 40

80% — Good

The frontend is well aligned with the problem statement: it appears to provide a command center, route optimization, shipment tracking, AI recommendations, alerts, and analytics. Where static evidence was limited, I relied on the FE capability map and retrieved snippets; I did not treat the Integration Map's backend-recognition gap as negative because it was explicitly marked unavailable. The main deductions are for incomplete direct evidence of loading/accessibility details and one flagged unmatched `/health` call.

Strengths

- Broad coverage of required logistics user journeys across planner dashboard, route optimization, AI recommendations, analytics, and public shipment tracking.
- Good evidence of user feedback and graceful degradation via alerting, polling resilience, and deterministic fallback behavior.

Areas to Improve

- Limited direct evidence of explicit loading states and robust accessibility implementation.
- Static analysis flagged at least one unmatched frontend API call (`/health`), so not all flows are fully validated end-to-end.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	UI covers the core user journeys required by the problem statement. Evidence: Capability Map / FE detail cites `frontend/src/pages/CommandCenter.tsx` with a live map, `Dashboard.tsx`, `Analytics.tsx`, and public `TrackShipment.tsx`; also `RecommendationCard.tsx` for accept/reject of AI suggestions...	MET	8/8
2	Frontend is functional and internally consistent (no obviously broken flows). Evidence: `frontend/package.json` includes runnable scripts (`"dev": "vite", "build": "tsc -b && vite build", "typecheck": "tsc --noEmit") and React/Router/Zustand dependencies. FE detail notes defined routes `/`, `/dashboard...	PARTIAL	6/8
3	Reasonable UX practices: loading states, error states, feedback to user. Evidence: FE detail states `src/services/realTimeService.ts` polls every 5 seconds and uses `Promise.allSettled` to handle partial API failures gracefully, pushing throttled connectivity alerts to the UI. It also says `AlertPanel...	PARTIAL	6/8
4	UI matches or closely reflects what is shown in the PPT screenshots. Evidence: PPT Slide 1/3 presents an "AI Logistics Command Center" with active simulation, multi-agent AI, route optimization, freight tracking, human-in-the-loop, live map tracking, ETAs, risk alerts, and KPI analytics. FE detail ...	MET	8/8
5	Responsiveness / accessibility considerations are present where relevant. Evidence: `frontend/package-lock.json` confirms `autoprefixer`, `tailwindcss`, and React-based SPA tooling, and FE detail notes the app shell in `App.tsx` has a collapsible sidebar plus reusable modal/panel components (`AddShipment...	PARTIAL	4/8

75% — Good

The code shows meaningful LLM integration in a multi-agent logistics workflow, especially for traffic/weather reasoning and explainable recommendations, which aligns well with the problem statement's generative-AI requirement. Static evidence did not detect a formal LLM SDK or vector store, but per instructions I did not treat NOT_DETECTED as a failure because the retrieved code clearly shows Gemini REST usage and tool-based retrieval. The main scoring deductions come from limited explicit prompt-injection handling and lightweight context management rather than from lack of relevance.

Strengths

- LLM use is tightly connected to operational recommendations, SLA impact explanation, and human-in-the-loop logistics decisions.
- Prompts and outputs are structured with explicit JSON schemas, fail-soft parsing, and governance validation against deterministic logic.

Areas to Improve

- No strong evidence of explicit prompt-injection defenses or sanitization for untrusted tool/user content.
- Context handling relies on live tool retrieval without evidence of richer memory, chunking, or benchmarked latency optimization.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	LLM is used purposefully to solve part of the core problem (not decorative). <i>Evidence: backend/app/services/llm_service.py: prompt asks Gemini to synthesize logistics recommendation context into JSON with fields like "sla_impact", "recommended_action", and "human_approval_required"; backend/app/main.py: "R..."</i>	MET	9/10
2	Prompts are structured, with clear instructions and/or few-shot guidance. <i>Evidence: backend/app/services/llm_service.py: explicit instruction "Return this exact JSON structure" followed by required schema fields; backend/app/llm/gemini_client.py docstring states "schema spelled out in the system prompt"...</i>	MET	8/10
3	Appropriate model/parameters chosen for the task (cost/latency/quality tradeoff). <i>Evidence: backend/app/services/llm_service.py uses "model='gemini-3.6-flash'"; backend/app/llm/gemini_client.py sets DEFAULT_MODEL to "gemini-flash-latest" and uses "temperature: float = 0.2"; same file describes it as "dependen..."</i>	MET	8/10
4	Output parsing / structured output handling is robust (JSON, function calling, etc.). <i>Evidence: backend/app/services/llm_service.py strips markdown fences from Gemini output before parsing; backend/app/llm/gemini_client.py says "Ask Gemini for a JSON response. Returns the parsed dict, or None on any failure" and us...</i>	MET	9/10
5	Prompt injection / unsafe input handling is considered. <i>Evidence: backend/app/llm/gemini_client.py docstring: "reason only from the given data"; backend/app/main.py and agent flow use MCP/tool-mediated live data plus deterministic fallback; backend/app/engines/governance.py checks align...</i>	PARTIAL	4/10
6	Context management (chunking, memory, retrieval) is appropriate to the use case. <i>Evidence: Capability map and backend/app/main.py describe agents using an MCP tool call -> Gemini reasoning pattern to fetch live traffic/weather context before analysis; backend/app/mcp_tools/server.py exposes geocoding, weather,...</i>	PARTIAL	7/10

90% — Good

This is a strong hackathon-grade AI engineering implementation centered on a coherent multi-agent pipeline, MCP-style tool use, governance, and human-in-the-loop execution that matches the logistics problem well. The main gap versus the success checklist is not architecture but evidence of formal AI evaluation/testing and privacy/bias controls. Where the rubric asked about RAG, I did not over-penalize the absence of vector-store-style document RAG because the problem statement is better served here by live retrieval from traffic/weather/route tools.

Strengths

- Clear multi-agent workflow with specialist agents, risk aggregation, deterministic decisioning, and governance/human approval.
- Strong resilience through deterministic fallback when Gemini is unavailable and partial-failure handling in real-time data refresh.

Areas to Improve

- Limited evidence of formal AI evaluation, latency/load observability, or benchmarking against the stated performance targets.
- Responsible AI evidence is stronger on transparency/governance than on bias awareness, anonymization, or privacy-compliance enforcement.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	Overall system architecture (agents, pipelines, tools) is coherent and justified. Evidence: backend/app/services/orchestrator.py: imports and sequences `analyze_traffic`, `analyze_weather`, `analyze_fleet`, `analyze_hub`, `analyze_coload`, `calculate_risk`, `make_decision`, and `validate_ai_reasoning`; backend/...	MET	8/9
2	RAG / retrieval pipeline (if used) is implemented soundly (chunking, indexing, retrieval). Evidence: backend/app/mcp_tools/server.py: MCP server exposes live tools like `geocode_location` and route/weather/traffic references; frontend/src/components/recommendations/RecommendationCard.tsx comments state agents "reason ov..."	PARTIAL	4/9
3	Tool/function calling or agentic orchestration is correctly wired end-to-end. Evidence: backend/app/mcp_tools/server.py: `@server.tool()` functions expose agent tools through MCP; backend/app/main.py startup text prints `AI agents: traffic_agent + weather_agent via MCP tools`; backend/app/services/orchestra...	MET	8/9
4	Evaluation, logging, or observability of AI behavior is present in some form. Evidence: backend/app/services/orchestrator.py prints `TRAFFIC:`, `WEATHER:`, `FLEET:`, `HUB:`, `COLOAD:`, `ROUTE:` intermediate outputs; capability map notes `POST /api/actions` and `GET /api/actions/audit` for action/audit loggi...	PARTIAL	7/9
5	System degrades gracefully on AI/model failures (fallbacks, retries). Evidence: backend/app/main.py startup text: `Gemini key not set — using deterministic fallback`; frontend/src/components/recommendations/RecommendationCard.tsx comment mentions `deterministic fallback when no API key is`; capabili...	MET	8/8

6 Engineering choices are appropriate for a hackathon time-box (pragmatic, not over-engineered).

MET

8/8

Evidence: Capability map: hybrid approach uses in-memory demo data (`app/data/shipments.py`, `app/data/vehicles.py`) alongside SQLAlchemy schema; frontend polls every 5 seconds via `realTimeService`; PPT Slide 3 labels it a `Worki...

7 Verify that responsible AI practices including transparency and bias awareness are demonstrated.

PARTIAL

6/8

Evidence: backend/app/engines/governance.py sets `human\_approval\_required` based on confidence/SLA/action alignment and returns `authoritative_source: "deterministic_decision_engine"; backend/app/services/governance_service.py co...

Security & HITL (SECURITY_HITL)

26 / 40

65% — Average

This submission is strongest on human-in-the-loop governance: the backend computes approval requirements, the decision engine flags high-risk cases, and the frontend supports accept/reject review of AI recommendations before operational action. Security hygiene is decent on secrets/configuration, but authentication/authorization is notably absent in the evidence, which matters because the product includes planner dashboards and shipment actions. The problem statement also explicitly calls for privacy compliance; since the code shows only limited visible privacy controls, I scored that partially rather than assuming compliance from the PPT alone.

Strengths

- Clear human-in-the-loop safety design with governance checks, approval flags, and operator accept/reject workflow.
- Good baseline secret hygiene: environment-variable configuration and no hardcoded secrets detected.

Areas to Improve

- No concrete authentication/authorization evidence for dashboards or operational endpoints.
- Limited explicit input sanitization and privacy-control evidence beyond basic UI constraints and low-PII demo data.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	Sensitive operations have human-in-the-loop confirmation where appropriate.	MET	8/8
Evidence: backend/app/services/decision_engine.py: vehicle breakdown returns recommendation with "Human intervention is required before continuing the shipment." and `approval_required=True`; backend/app/main.py: recommendation re...			
2	Basic security hygiene: no secrets committed, input sanitization present.	PARTIAL	5/8
Evidence: STATIC EVIDENCE: "No hardcoded secrets/API keys in source: PASS" and "Configuration via environment variables: PASS"; backend/app/services/llm_service.py uses `os.getenv("GEMINI_API_KEY")`; backend/app/main.py loads env ...			
3	Authentication/authorization is present where the use case requires it.	NOT_MET	1/8
Evidence: STATIC EVIDENCE: "Known auth/authorization library or pattern used: NOT_DETECTED". Capability Map notes a public-facing `TrackShipment.tsx` and planner dashboards, but no login, roles, token checks, or protected-route/ba...			

4 AI outputs that affect real-world actions have a review/override mechanism.

MET

8/8

Evidence: Capability Map: `RecommendationCard.tsx` lets users accept or reject AI suggestions; backend `GovernanceEngine` validates AI-generated recommendations and sets `human\_approval\_required`; backend/app/main.py returns both ...

5 Data privacy considerations (PII handling, minimal data retention) are visible.

PARTIAL

4/8

Evidence: Problem success checklist explicitly requires "Compliance with data privacy regulations and use of high-quality, anonymized data." The code mainly uses operational shipment/vehicle/hub fields and in-memory demo data in b...

MVP & Outcome (MVP)

44 / 50

Excellent

The submission is a credible MVP with substantial code evidence for a logistics command center, multi-agent reasoning, route optimization, tracking, and explainable recommendations. I anchored scoring to the problem statement: dashboard, multi-agent flow, adaptive routing, and customer/planner value are well represented, while missing proof for checklist items like performance targets, privacy compliance, and robust integration/load testing prevented top scores. The integration map explicitly lacked recognized backend detection, so I treated that as unavailable rather than negative and relied on the retrieved code and capability map.

Strengths

- Working full-stack prototype with live simulation, route optimization outputs, and visible frontend flows for planners and trackers.
- Strong alignment with the core challenge through multi-agent orchestration, explainable AI recommendations, governance, and real-time dashboard-style UX.

Areas to Improve

- Limited hard evidence for success-checklist metrics such as 20% delivery-time reduction, sub-3 second recalculation latency, and privacy/compliance handling.
- Backend/frontend integration could not be fully statically corroborated, and automated testing evidence appears present but not extensive for integration/load validation.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	The submission is a working, demonstrable MVP (not just slides/mockups).	MET	8/9
Evidence: backend/app/main.py: FastAPI app is instantiated as a "Simple Working Version" and contains live demo simulation logic; retrieved snippet shows shipment progress is updated over time in `_simulate_shipments()` and route ...			
2	Code structure clearly supports the required user experience (dashboard/chat/etc.).	MET	8/9
Evidence: Capability map cites frontend pages `CommandCenter.tsx`, `Dashboard.tsx`, `WarehouseDashboard.tsx`, `Analytics.tsx`, and `TrackShipment.tsx`; `GoogleMapComponent.tsx` and `RoutePreviewMap.tsx` support map-based tracking ...			
3	Core happy-path of the use case can be traced end-to-end through the code.	MET	7/8
Evidence: backend/app/main.py route optimization flow geocodes origin/destination and calls routing tools (`call_tool("geocode_location"...)`), then returns route and risk data. frontend/src/ty...			

4	PPT screenshots/demo descriptions are corroborated by actual code artifacts.	MET	7/8
Evidence: Slide 3 claims "Traffic Agent, Weather Agent, Fleet Agent, Hub Agent, Co-load Agent, Risk Agent," decision engine, explanation, governance, command center, and human approval. Code corroboration appears in `backend/app/a...			
5	Scope is realistic and mostly complete for the hackathon time constraints.	MET	7/8
Evidence: backend/app/main.py is explicitly described as a "Simple Working Version" scoped to "Nagpur city and its nearby localities." Slides 1 and 3 label it a "Working Prototype / Simulated Data." The implementation uses in-memo...			
6	Verify that outputs provide value to end users.	MET	7/8
Evidence: backend/app/main.py returns actionable outputs including `recommended_route`, `traffic_warnings`, `weather_warnings`, `recommended_vehicle`, `risk_level`, and ETA/delay estimates. frontend/src/components/recommendations/...			

Presentation & Innovation (PRESENTATION)
25 / 37



The presentation is strong on clarity of problem framing, system architecture, and alignment with the implemented multi-agent logistics prototype. I weighted the problem statement heavily: since the challenge asks for measurable logistics improvements, I penalized the lack of quantified business impact and absent explicit roadmap despite otherwise solid storytelling and technical coherence. Code and capability evidence support that the deck is not merely aspirational, but the PPT excerpts are more architecture-centric than demo- and metrics-centric.

Strengths

- Clear problem-to-solution narrative centered on multi-agent logistics orchestration and human-in-the-loop governance.
- PPT claims are largely backed by code: live map/dashboard surfaces, route optimization, tracking flows, explainable AI, and governance logic.

Areas to Improve

- Business impact is not quantified against the success checklist metrics such as 20% delivery-time reduction or sub-3s recalculation.
- Future roadmap is not explicitly presented in the provided slides, making long-term realism harder to assess.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	PPT clearly explains problem, approach, and impact.	MET	5/6
Evidence: PPT Slide 2 states the problem directly: "Difficulty optimizing delivery routes dynamically against real-time conditions (traffic, weather, vehicle availability)." Slide 3 lays out the approach with "DATA INPUTS," "MULTI...			
2	Innovation / originality of the approach is evident.	MET	5/6
Evidence: PPT Slide 3 presents a nontrivial pipeline: specialized agents, deterministic decision engine, Gemini explanation layer, and governance/human approval. Code evidence shows originality in backend/app/engines/sla_engine.py...			
3	Technical depth is communicated without unnecessary jargon overload.	PARTIAL	4/5
Evidence: PPT Slide 3 communicates architecture in plain blocks like "Decision Engine," "Explanation," and "Governance" rather than dense ML terminology. This matches the code's readable comments, e.g., frontend/vite.config.ts: "L...			

4	Demo narrative (screens, flow) matches the actual implementation.	PARTIAL	4/5
Evidence: PPT Slide 3 promises a "COMMAND CENTER" with "Live map tracking Delivery ETAs Risk alerts KPI Analytics." The codebase contains matching screens/components: frontend/src/pages/WarehouseDashboard.tsx includes "Co-load opp...			
5	Overall polish and storytelling of the submission.	PARTIAL	4/5
Evidence: Slides 1-3 show a coherent narrative progression: "THE PROBLEM" -> "THE SYSTEM" -> "AI ARCHITECTURE," with consistent branding and concise sectioning. The frontend also reflects product polish via components like Persona...			
6	Verify that business impact is quantified where possible.	NOT_MET	1/5
Evidence: No clear quantified business impact matching the success checklist was found in the visible PPT excerpts or retrieved code. While the app includes KPI/analytics surfaces and backend/app/services/decision_engine.py refere...			
7	Verify that future roadmap is realistic.	PARTIAL	2/5
Evidence: No explicit roadmap slide is included in the provided PPT excerpts. The code and comments suggest a realistic prototype scope—e.g., backend/app/main.py and frontend/src/utlis/geo.ts openly constrain the demo to "Nagpur c...			

Innovation & Business Value (INNOVATION)

9 / 10

ent

The solution is genuinely innovative for this problem because it combines multi-agent analysis, tool-augmented AI, deterministic decisioning, governance, and human approval in a logistics command-center workflow. I weighted the problem statement heavily: this helped the team on AI creativity and architectural relevance, but I held back full business-value points because the checklist asks for concrete performance proof and the evidence mainly shows KPI reporting surfaces rather than validated outcome measurements. The Integration Map explicitly lacked recognized backend signals, so I did not penalize that and instead relied on the Capability Map and retrieved code snippets as instructed.

Strengths

- Strong multi-agent + governance design tailored to dynamic logistics decisions rather than generic chat-style AI.
- Clear operational and customer value surfaces through KPIs, alerts, ETAs, tracking, and explainable recommendations.

Areas to Improve

- Business impact appears instrumented in dashboards, but hard evidence for the required 20% improvement and latency target is not shown.
- Prototype still depends partly on in-memory/demo data, which limits proof of production-scale execution despite a scalable design.

RULE EVALUATIONS

#	Rule	Verdict	Score
1	AI techniques are combined creatively to solve the problem.	MET	3/3
Evidence: PPT Slide 3 describes 'Traffic Agent, Weather Agent, Fleet Agent, Hub Agent, Co-load Agent, Risk Agent' feeding a 'Decision Engine' with Gemini used for 'human-readable explanations'; Capability Map states `app/orchestra...			

2

Solution provides measurable business or customer value beyond basic functionality.

PARTIAL

2/3

Evidence: frontend/src/pages/Dashboard.tsx defines KPI cards for 'SLA Compliance', 'Delivery Time Reduction', and 'Cost Savings'; frontend/src/services/realTimeService.ts pushes critical alerts such as 'Critical risk shipments' an...

3

Architecture demonstrates scalable innovation.

MET

2/2

Evidence: frontend/src/services/api.ts centralizes typed API access; frontend/src/services/realTimeService.ts uses polling with `Promise.allSettled` to tolerate partial failures; backend/app/engines/sla_engine.py isolates SLA logi...

4

Solution demonstrates thoughtful AI engineering practices beyond simple API integration.

MET

2/2

Evidence: Capability Map notes `GovernanceEngine` validates AI recommendations against deterministic business rules and sets `human\_approval\_required`; `TrafficAgent` and `WeatherAgent` fetch live tool data first, then analyze it ...

Static Analysis Checks

Check	Verdict
No hardcoded secrets/API keys in source No hardcoded-secret-like patterns matched.	PASS
Configuration via environment variables ai-logistics-command-center/frontend/src/services/api.ts, ai-logistics-command-center/backend/app/main.py, ai-logistics-command-center/backend/app/services/llm_service.py	PASS
Automated test files present test_llm.py, test_routing.py, test_weather.py	PASS
CI/CD pipeline configured	NOT_DETECTED
Known input-validation library used	NOT_DETECTED
Known auth/authorization library or pattern used	NOT_DETECTED
Recognized LLM SDK/framework in use	NOT_DETECTED
Recognized vector store / RAG library in use	NOT_DETECTED
Agentic/tool-calling orchestration pattern detected	NOT_DETECTED
Recognized database driver/ORM in use	NOT_DETECTED
try/except/catch present in backend-ish code ai-logistics-command-center/backend/app/llm/gemini_client.py, ai-logistics-command-center/backend/app/mcp_tools/client.py, ai-logistics-command-center/frontend/src/components/map/Go	PASS

Problem Summary

****Summary:**** The challenge is to develop a generative AI-driven multi-agent system for dynamic route optimization and freight tracking in logistics, addressing inefficiencies caused by real-time conditions. The solution should enhance delivery efficiency, reduce costs, and improve customer satisfaction.

Success Checklist:

- Multi-agent architecture for traffic/weather data ingestion, vehicle tracking, and route optimization.
- Real-time dashboard for logistics planners and customers with delivery ETAs and alerts.
- Generative AI algorithms for adaptive routing and scenario simulations.
- API integrations for seamless data flow from traffic, weather, and GPS sources.
- Performance metrics showing at least a 20% reduction in delivery times and sub-3 second route recalculation latency.
- Comprehensive documentation for deployment and user operations.
- Automated testing for integration and load.
- Compliance with data privacy regulations and use of high-quality, anonymized data.

Cross-Validation Report

Cross-validation report

1) Cross-category inconsistencies

- Backend framework certainty is overstated in some places.
- The Capability Map says the backend is “built with FastAPI.”
- But the Integration Map is explicit ground truth for static wiring and says no routes were detected and no recognized backend framework imports were found.
- The BACK result still says “credible FastAPI-based logistics prototype.” That may be true from retrieved snippets, but it is too certain relative to the shared evidence. It should be phrased more cautiously as “Python backend described as FastAPI in capability evidence, but static integration could not verify framework/routes.”
- FE!“BE integration is assumed more strongly than the wiring evidence supports.
- FRONT, MVP, and PRESENTATION all speak as if the UI is wired to backend features broadly.
- Ground truth integration shows only 1 frontend API call found and 0 matched FE!“BE pairs, with
- Because the Integration Map itself says backend detection is unavailable, this is not proof the app lacks integrations, but it does mean categories should avoid strong claims of verified end-to-end API wiring. They should say “feature evidence exists in code/capability map, but static FE!“BE linkage was not verified.”
- HITL execution path is described more concretely than proven.
- Capability Map claims frontend actions trigger `POST /recommendations/{id}/accept`.
- Integration Map found zero backend routes.
- Again, not a contradiction if route detection is unavailable, but any category presenting this as verified end-to-end behavior is overstating confidence.

2) PPT vs Code gaps

- “6 AGENTS ACTIVE” in the PPT is not fully backed by the shared code narrative.
- PPT names Traffic, Weather, Fleet, Hub, Co-load, Risk.
- Capability/code narrative explicitly supports Traffic, Weather, Fleet, Hub, plus orchestrator/decision/governance.
- Co-load Agent and Risk Agent are not clearly substantiated in the provided code summary. Risk may be partially implied by `risk\_assessments` table and governance/decision logic, but that is weaker than explicit agent evidence.
- This is the clearest PPT-over-code gap.
- “Gemini generates human-readable explanations” appears broadly supported, but “separation of decision logic and narration” is more architectural than directly evidenced.
- DecisionEngine + governance + LLM narration is plausible from the narrative, but categories should be careful not to overclaim strict separation unless code snippets showed it directly.
- Human-in-the-loop in the PPT is directionally supported by both frontend/backend narrative, but specific API execution claims are not verified by static integration.
- So HITL exists as a concept and likely implementation, but “fully wired accept/reject workflow” is not statically backed.
- Persistent database story is weaker than a full production DB-backed system.
- Code evidence supports ORM models/tables.
- Capability Map also says runtime primarily uses in-memory data for demos.
- Any PPT implication of fully persistent operational data should be treated cautiously.

3) Likely score miscalibration

- BACK 41/50 may be a bit high given the verification limits.
- Not because the backend is weak, but because framework/routes/API wiring were not statically verified and persistence is partly demo/in-memory.
- A slightly lower score or lower confidence would better match the evidence posture.
- FRONT 32/40 may be somewhat high if interpreted as implemented-and-wired rather than UI capability present.
- The UI feature surface sounds strong, but FE!“BE verification is extremely thin in the integrati

- If the rubric rewards end-to-end product readiness, this should be tempered slightly or confidence reduced.
- LLM 45/60 and AI_ENG 49/60 look broadly reasonable.
- They appropriately credit tool-augmented multi-agent reasoning without forcing vector-store RAG.
- No major inconsistency there.
- SECURITY_HITL 26/40 looks calibrated.
- It gives credit for governance/HITL but penalizes missing auth/privacy evidence, which is consistent with the shared record.
- PRESENTATION 25/37 looks fair, possibly even generous.
- The deck is coherent, but some architecture claims—especially agent count—outrun the clearly substantiated code evidence.

Bottom line

- Main evidence risk: several categories speak as if backend routes and FE!"BE execution are known, but wiring is unavailable/not verified.
- Main PPT-code gap: the slide claim of "6 agents active" is not fully supported by the code narrative provided.
- Most likely over-scored areas: BACK and possibly FRONT unless their scores are understood as prototype capability rather than verified end-to-end implementation.

Final Evaluation Report

Evaluation Report

Overall Score: 293 / 377

Confidence: Medium

Category Breakdown

- Problem Statement (PS):** 22/30
- Backend (BACK):** 41/50
- Frontend (FRONT):** 32/40
- LLM Usage (LLM):** 45/60
- AI Engineering (AI_ENG):** 49/60
- Security & HITL (SECURITY_HITL):** 26/40
- MVP & Outcome (MVP):** 44/50
- Presentation & Innovation (PRESENTATION):** 25/37
- Innovation & Business Value (INNOVATION):** 9/10

Detailed Analysis

Strengths

The submission is a strong logistics-focused AI prototype with a coherent command-center narrative carried through slides, backend services, and frontend surfaces. Its biggest strength is the multi-agent workflow: traffic, weather, fleet, hub, risk/governance-style analysis, deterministic decisioning, and human-readable explanations are all clearly aimed at the actual operational problem rather than generic chatbot behavior. The system also shows good hackathon-grade engineering judgment through fallback behavior when Gemini is unavailable, live-tool retrieval for route/weather/traffic context, and human-in-the-loop approval for higher-risk actions.

On product scope, the team appears to cover the right user journeys: planner dashboards, live maps, ETA/risk visibility, route optimization, alerts, recommendations, and public shipment tracking. The presentation is also well aligned with the codebase at a high level: the problem framing, architecture, and command-center workflow are credible and mostly supported by implementation evidence.

Weaknesses & Gaps

The main gap is evidence quality rather than concept quality. The project repeatedly claims operational impact, but there is no hard proof for key success metrics such as 20% delivery-time improvement or sub-3-second recalculation latency. Formal AI evaluation is also limited: there is little visible benchmarking, latency/load observability, regression testing, or systematic validation of model quality.

Security and responsible-AI controls are only partial. Secret handling and environment-based configuration are good, but there is no concrete evidence of authentication/authorization for operational dashboards or shipment actions, and privacy/compliance measures are not strongly demonstrated. Persistence is also mixed: ORM/database structures exist, but runtime behavior appears to rely substantially on in-memory or demo data, which weakens claims of production readiness.

PPT vs Code Discrepancies

The clearest discrepancy is the PPT claim of "6 agents active." Traffic, weather, fleet, and hub are reasonably supported by the code narrative, but co-load and risk as distinct implemented agents are not equally well substantiated in the provided evidence. Risk is partly supported through risk calculation/governance logic, but that is weaker than explicit agent implementation.

There is also a verification gap between UI/backend claims and static integration evidence. The deck and some category results imply broad end-to-end functionality, but static FE/BE wiring was not verified; only an unmatched "/health" human-in-the-loop is directionally supported in both PPT and code, but a fully verified accept/reject execution path was not statically established.

Innovation Notes

This is a genuinely thoughtful application of AI to logistics operations. The novelty is not just "using an LLM," but combining tool-augmented specialist analysis, deterministic decision logic, governance, and operator approval into a workflow that fits dispatch/exception-management realities. The solution also creates visible business-value surfaces—ETA transparency, SLA impact explanation, route/risk recommendations, and customer tracking—even if measured outcome evidence is still incomplete.

Rule-level Evidence

- ****Problem Statement (PS)**:**
- ``backend/app/services/orchestrator.py`` computes dynamic ETA using route duration plus traffic/weather/hub/co-load delays.
- ``frontend/src/components/map/GoogleMapComponent.tsx`` shows live map tracking with vehicles, shipment markers, routes, and hubs.
- PPT Slide 3 presents the "AI Logistics Command Center" with live map tracking, ETAs, risk alerts, and analytics.
- ****Backend (BACK)**:**
- ``backend/app/api/simulation.py`` and ``backend/app/routes/recommendations.py`` show structured router usage and domain-specific API grouping.
- ``backend/app/database.py``, ``backend/app/models/shipment.py``, ``route.py``, and ``audit.py`` show ORM-backed persistence intent.
- However, ``backend/app/main.py`` and ``app/data/shipments.py`` evidence demo/in-memory state in live flows, weakening persistence maturity.
- Cross-validation note: framework/routes are described via retrieved snippets, but static integration did not verify backend framework recognition.
- ****Frontend (FRONT)**:**
- ``frontend/src/pages/CommandCenter.tsx``, ``Dashboard.tsx``, ``Analytics.tsx``, ``TrackShipment.tsx``, and ``RouteOptimizer.tsx`` indicate broad logistics UX coverage.
- ``frontend/src/components/recommendations/RecommendationCard.tsx`` supports recommendation review and reasoning display.
- ``frontend/package.json`` includes runnable scripts and core app dependencies.
- Integration caveat: ``frontend/src/services/api.ts`` contains a detected ``${rootURL}/health`` call that was not matched to a verified backend route.
- ****LLM Usage (LLM)**:**
- ``backend/app/services/llm_service.py`` uses structured prompts requiring exact JSON outputs with fields such as ``sla_impact``, ``recommended_action``, and ``human_approval_required``.
- ``backend/app/llm/gemini_client.py`` documents reasoning from provided data and schema-constrained behavior.
- PPT Slide 3 says Gemini generates human-readable explanations.
- Static evidence did not detect a formal vector store or richer long-context system.
- ****AI Engineering (AI_ENG)**:**
- ``backend/app/services/orchestrator.py`` sequences multiple analysis modules plus risk calculation, decisioning, and AI-reasoning validation.
- ``backend/app/mcp_tools/server.py`` exposes live retrieval/tool interfaces such as geocoding and route/weather/traffic references.
- The architecture shows deterministic fallback and partial-failure handling, but little formal evaluation/benchmark evidence.
- ****Security & HITL (SECURITY_HITL)**:**
- ``backend/app/services/decision_engine.py`` explicitly marks certain recommendations as requiring human intervention.
- ``backend/app/main.py`` returns ``approval_required`` driven by governance logic.
- Static evidence: no hardcoded secrets detected; environment-based configuration present.
- Gaps remain around auth/authz, privacy controls, and robust input validation ("Known input-validation library used: NOT_DETECTED").
- ****MVP & Outcome (MVP)**:**
- ``backend/app/main.py`` is described as a working prototype with simulation logic, and ``_simulate_shipments()`` updates shipment progress over time.
- Route optimization outputs include concrete operational fields like ``distance_km``, ``estimated_duration_hours``, ``risk_score``, and ``recommended_route``.
- ``frontend/src/hooks/useShipments.ts`` and real-time services support a live-feeling dashboard experience.
- Slide 1 explicitly claims "STABLE PROTOTYPE" and "ACTIVE SIMULATION."
- ****Presentation & Innovation (PRESENTATION)**:**
- PPT Slide 2 clearly states the logistics problem; Slide 3 explains the multi-agent architecture and governance flow.
- Code in ``backend/app/services/directions_service.py`` and ``backend/app/services/llm_service.py`` supports the architectural story.
- The deck is coherent, but it is more architecture-centric than metrics-centric, and some claims exceed what is directly substantiated.
- ****Innovation & Business Value (INNOVATION)**:**
- The combination of specialist agents, tool-based retrieval, deterministic decisioning, governance, and explainability is well tailored to dynamic logistics operations.

- `frontend/src/pages/Dashboard.tsx` includes KPI surfaces such as SLA compliance, delivery time reduction, and cost savings.
- `frontend/src/pages/TrackShipment.tsx` adds customer-facing value.
- However, KPI presence is not the same as validated business impact.

Recommended Judge Questions

1. Your slides claim **6 active agents**. Which exact files implement each of the six, especially the **Co-load Agent** and **Risk Agent**?
2. What evidence do you have for the stated operational outcomes—especially **delivery-time reduction** and **recalculation latency**—beyond dashboard KPI displays?
3. How is the system secured for real operational use: do you have **authentication, authorization, and audit controls** for planner actions and recommendation approvals?
4. Which parts of the workflow are truly **database-backed** today, and which are still using **in-memory/demo state**?
5. How do you protect the LLM pipeline from **prompt injection or untrusted tool/user input**, especially when using live external data?
6. If Gemini is unavailable or produces low-confidence output, what exact fallback path executes and how does that affect operator experience?

Evaluation Confidence & Coverage

- Code ingested: 124 file(s) across 3 layer(s) (FE/BE/DB/OTHER).
- Backend framework recognition: NONE recognized (63 backend-layer file(s) scanned). Static route/integration analysis found no known-framework signal - treat Integration Map findings as unavailable, not negative.
- 4/11 deterministic checks found concrete evidence (PASS/FAIL); 7 had no signal (NOT_DETECTED = absence of evidence, not evidence of absence).
- RAG retrieval: BM25 + embeddings + cross-encoder reranker.
- Knowledge graph: 261 nodes, 354 edges - Graph-RAG context available.